Kuku Promo Finder

Problem 1 — Solution overview

Problem

Identify 20–60-second promos from a microdrama screenplay. Each clip should have a hook in the first three seconds, make sense to a new viewer and end on an unresolved moment. Return exact segment timestamps and explain which promos to post first.

Approach

The workflow is parsing → candidate generation → deduplication → ranking. Python handles source records, timestamps and validation; Gemini handles story interpretation and editorial scoring. Episodes are searched together, without a promo quota for each episode.

Parsing

The parser reads scenes, dialogue, actions, camera directions and sound cues. It preserves original timestamps and assigns a source ID to each entry. Overlapping audiovisual entries stay separate. Malformed records raise errors; timing inconsistencies are flagged rather than corrected silently.

Candidate generation

Gemini 3.7 receives ten episodes per window, advancing by five episodes. Each response can propose up to eight promos, including continuous clips or ordered segments across episodes. The model selects source IDs; code retrieves their timestamps and checks duration and boundaries. For example, a segment from 00:12 to 00:47 contains 35 seconds of original footage.

Deduplication

Cuts are grouped when corresponding segments use the same episodes and playback order and overlap by at least 80% of the shorter segment. The first occurrence is retained. The full 50-episode pack produced 51 technically valid candidates and 35 groups; duplicate alternatives remain saved for reference.

Ranking

The current judge scores only the retained candidates, using their actual screenplay excerpts without generator explanations or full-story summaries. It gives evidence-backed scores from 0–4 for hook, cliffhanger, clarity and escalation.

Code weights these scores at 30%, 30%, 25% and 15%. A score below 2 on hook, cliffhanger or clarity flags a promo; it is retained below recommended candidates. A final comparison orders the leading shortlist. JSON and CSV exports contain ranks, timestamps, explanations and individual scores.

Run note: the existing saved ranking predates this simplification and also assessed duplicate alternatives.

Limitations and improvements

Free-tier quotas limited model comparisons. The run used medium thinking and a saved request budget. With more capacity, stronger models and higher thinking levels could be tested against human-reviewed examples to measure whether they improve selections.

The judge reads text rather than finished video, may share the generator’s biases, and cannot establish exact word timing. An editor should review cut points and top selections. Scores are editorial judgments, not advertising results.

AI use and checks

Codex assisted planning, coding and documentation; Gemini proposed and scored promos. An early token limit caused truncated JSON, so finish-status checks now reject incomplete responses. Automated tests cover parsing, timing, duplicates, judge evidence and exports.